

Gestational Diabetes Dataset

Source: <https://www.kaggle.com/datasets/rasooljader/gestational-diabetes>

```
[23]: import pandas as pd
df=pd.read_csv('/content/Gestational Diabetes.csv')
df
```

```
[23]:
```

| | Age | Pregnancy No | Weight | Height | BMI | Heredity | Prediction |
|------|------|--------------|--------|--------|------|----------|------------|
| 0 | 17.0 | 1.0 | 48.0 | 165.0 | 17.6 | 0 | 0 |
| 1 | 17.0 | 1.0 | 49.0 | 145.0 | 23.3 | 0 | 0 |
| 2 | 17.0 | 1.0 | 50.0 | 140.0 | 25.5 | 0 | 0 |
| 3 | 17.0 | 1.0 | 50.0 | 145.0 | 23.8 | 0 | 0 |
| 4 | 17.0 | 1.0 | 49.0 | 146.0 | 23.0 | 0 | 0 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 1007 | 35.0 | 3.0 | 89.0 | 159.0 | 35.2 | 1 | 1 |
| 1008 | 41.0 | 4.0 | 87.0 | 165.0 | 32.0 | 0 | 0 |
| 1009 | 34.0 | 2.0 | 67.0 | 160.0 | 26.2 | 1 | 0 |
| 1010 | 33.0 | 3.0 | 65.0 | 167.0 | 23.3 | 0 | 1 |
| 1011 | 28.0 | 2.0 | 68.0 | 156.0 | 27.9 | 0 | 0 |

[1012 rows x 7 columns]

```
[24]: df.head()
```

```
[24]:
```

| | Age | Pregnancy No | Weight | Height | BMI | Heredity | Prediction |
|---|------|--------------|--------|--------|------|----------|------------|
| 0 | 17.0 | 1.0 | 48.0 | 165.0 | 17.6 | 0 | 0 |
| 1 | 17.0 | 1.0 | 49.0 | 145.0 | 23.3 | 0 | 0 |
| 2 | 17.0 | 1.0 | 50.0 | 140.0 | 25.5 | 0 | 0 |
| 3 | 17.0 | 1.0 | 50.0 | 145.0 | 23.8 | 0 | 0 |
| 4 | 17.0 | 1.0 | 49.0 | 146.0 | 23.0 | 0 | 0 |

Exercise 1: Analyzing a Health-Related Dataset

```
[25]: # Calculate basic descriptive statistics
print("Mean:\n", df.mean())
```

```

print("\nMedian:\n", df.median())
print("\nMode:\n", df.mode().iloc[0])
print("\nStandard Deviation:\n", df.std())
print("\nVariance:\n", df.var())
# Additional descriptive statistics
print("\nRange:\n", df.max() - df.min())
print("\nSkewness:\n", df.skew())
print("\nKurtosis:\n", df.kurt())

```

Mean:

| | |
|--------------|------------|
| Age | 30.357708 |
| Pregnancy No | 2.517787 |
| Weight | 72.988142 |
| Height | 158.154150 |
| BMI | 29.256818 |
| Heredity | 0.226285 |
| Prediction | 0.214427 |

dtype: float64

Median:

| | |
|--------------|-------|
| Age | 30.0 |
| Pregnancy No | 2.0 |
| Weight | 70.0 |
| Height | 158.0 |
| BMI | 28.9 |
| Heredity | 0.0 |
| Prediction | 0.0 |

dtype: float64

Mode:

| | |
|--------------|-------|
| Age | 23.0 |
| Pregnancy No | 2.0 |
| Weight | 65.0 |
| Height | 160.0 |
| BMI | 31.6 |
| Heredity | 0.0 |
| Prediction | 0.0 |

Name: 0, dtype: float64

Standard Deviation:

| | |
|--------------|-----------|
| Age | 7.022960 |
| Pregnancy No | 1.530656 |
| Weight | 12.552712 |
| Height | 7.367421 |
| BMI | 5.092321 |
| Heredity | 0.418632 |
| Prediction | 0.410627 |

dtype: float64

Variance:

| | |
|--------------|------------|
| Age | 49.321968 |
| Pregnancy No | 2.342908 |
| Weight | 157.570581 |
| Height | 54.278885 |
| BMI | 25.931734 |
| Heredity | 0.175253 |
| Prediction | 0.168615 |

dtype: float64

Range:

| | |
|--------------|------|
| Age | 29.0 |
| Pregnancy No | 8.0 |
| Weight | 83.0 |
| Height | 61.0 |
| BMI | 38.3 |
| Heredity | 1.0 |
| Prediction | 1.0 |

dtype: float64

Skewness:

| | |
|--------------|----------|
| Age | 0.161460 |
| Pregnancy No | 1.256361 |
| Weight | 0.748643 |
| Height | 0.381251 |
| BMI | 0.665234 |
| Heredity | 1.310254 |
| Prediction | 1.393667 |

dtype: float64

Kurtosis:

| | |
|--------------|-----------|
| Age | -0.861067 |
| Pregnancy No | 1.534562 |
| Weight | 1.099075 |
| Height | 0.515864 |
| BMI | 1.073525 |
| Heredity | -0.283798 |
| Prediction | -0.057812 |

dtype: float64

```
[26]: import scipy.stats as stats

# Select the feature for analysis
feature = 'BMI' # Replace with your desired feature
```

```

# Define the hypothesized value
hypothesized_value = 29.42699440206 # Replace with your hypothesized value

# Extract the data for the selected feature
data = df[feature]

# Perform a one-sample t-test
t_stat, p_value = stats.ttest_1samp(data, hypothesized_value)

# Print the results
print(f"T-statistic: {t_stat}")
print(f"P-value: {p_value}")

# Determine if the hypothesis should be rejected
alpha = 0.05 # Significance level
if p_value < alpha:
    print("Reject the null hypothesis. The average BMI is significantly
    ↪different from the hypothesized value.")
else:
    print("The null hypothesis is accepted. There is no significant difference
    ↪between the average BMI and the hypothesized value.")

# Calculate the 95% confidence interval
confidence_level = 0.95
degrees_of_freedom = len(data) - 1
confidence_interval = stats.t.interval(confidence_level, degrees_of_freedom,
    ↪loc=data.mean(), scale=stats.sem(data))

# Print the confidence interval
print(f"95% Confidence Interval for the mean of {feature}:
    ↪{confidence_interval}")

```

T-statistic: -1.063098123519172

P-value: 0.2879913138277411

The null hypothesis is accepted. There is no significant difference between the average BMI and the hypothesized value.

95% Confidence Interval for the mean of BMI: (28.942699440206496, 29.57093692342987)

Exercise 2: Exploring Regression Analysis on a New Dataset

```

[27]: import statsmodels.formula.api as smf

# Define the regression model
model = smf.ols('Prediction ~ Age', data=df)

```

```
# Fit the model
results = model.fit()

# Print the summary of the regression results
print(results.summary())
```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          Prediction    R-squared:                  0.116
Model:                  OLS          Adj. R-squared:             0.115
Method:                 Least Squares  F-statistic:               132.6
Date:                  Thu, 05 Sep 2024  Prob (F-statistic):       6.32e-29
Time:                  14:35:13       Log-Likelihood:            -472.28
No. Observations:      1012          AIC:                      948.6
Df Residuals:          1010          BIC:                      958.4
Df Model:               1
Covariance Type:       nonrobust
=====

```

| | coef | std err | t | P> t | [0.025 | 0.975] |
|-----------|---------|---------|--------|-------|--------|--------|
| Intercept | -0.3903 | 0.054 | -7.243 | 0.000 | -0.496 | -0.285 |
| Age | 0.0199 | 0.002 | 11.517 | 0.000 | 0.017 | 0.023 |

```

=====
Omnibus:                 151.139    Durbin-Watson:              1.737
Prob(Omnibus):           0.000     Jarque-Bera (JB):           225.486
Skew:                    1.156     Prob(JB):                   1.09e-49
Kurtosis:                 3.001     Cond. No.                   138.
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

1 Explanation of the values:

2 R-squared:

This value indicates the proportion of variance in the dependent variable (Prediction) that is explained by the independent variable (Age).

A higher R-squared value suggests a stronger relationship between the variables.

3 Adjusted R-squared:

This is a modified version of R-squared that accounts for the number of predictors in the model. It helps to prevent overfitting.

4 Coefficients:

These values represent the estimated change in the dependent variable (Prediction) for a one-unit change in the independent variable (Age).

The intercept represents the predicted value of the dependent variable when the independent variable is zero.

5 P-values:

These values indicate the statistical significance of the coefficients. A p-value less than the significance level (usually 0.05) suggests that the coefficient is statistically significant.

6 Standard Errors:

These values indicate the variability of the estimated coefficients.

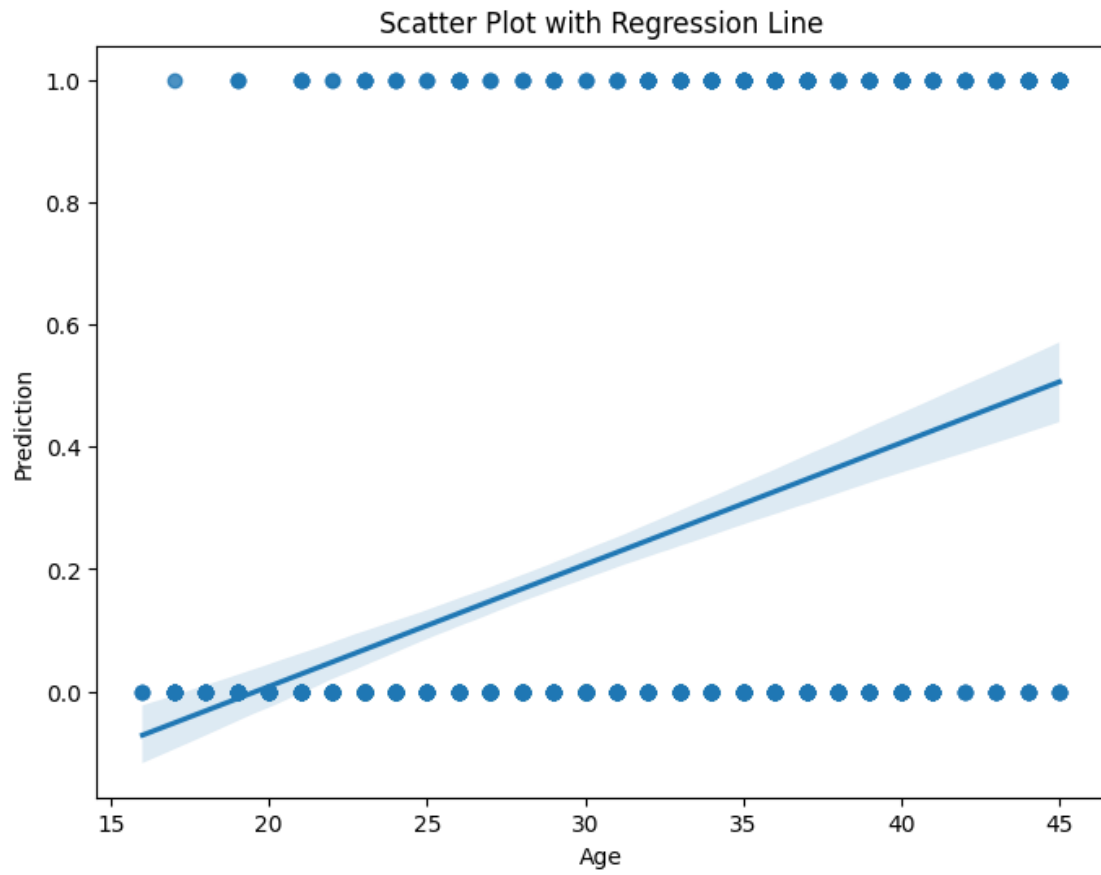
7 F-statistic and p-value:


These values test the overall significance of the regression model. A low p-value indicates that the model is statistically significant.

```
[28]: import matplotlib.pyplot as plt
import seaborn as sns

# Scatter plot with regression line
plt.figure(figsize=(8, 6))
sns.regplot(x='Age', y='Prediction', data=df)
plt.title('Scatter Plot with Regression Line')
plt.xlabel('Age')
plt.ylabel('Prediction')
plt.show()

# Histogram of Age
plt.figure(figsize=(8, 6))
sns.histplot(df['Age'], bins=10)
plt.title('Histogram of Age')
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.show()
```





[]: